AV ADR Vault

Sign in to your vault

Access your team's architectural decisions.

Email

mira@acme.dev

Password Forgot?

••••••••

Sign in

New to ADR Vault? Create an account

AV ADR Vault

Create your account

Start capturing architectural decisions today.

Display name

Mira Chen

Work email

mira@acme.dev

Password

••••••••

Minimum 8 characters.

Create account

Already have an account? Sign in

All ADRs

Every architectural decision captured by your team.

AI Brain Dump | New ADR

All 15 | Draft 1 | Proposed 1 | Accepted 12 | Deprecated 0 | Superseded 1

Search title, summary, or ADR-###

Sort: Last updated

ADR	TITLE	STATUS	CREATION	AUTHOR	UPDATED
ADR-001	**Authentication & authorisation flow** Session handling and auth mutations communicate directly with Supabase Auth.	Accepted	Manual	MC Mira	19d ago
ADR-002	**Tiptap as the WYSIWYG editor** Adopted for ADR authoring; JSON document tree stored as JSONB.	Accepted	AI · Edited	JA Jon	3w ago
ADR-003	**pgvector for semantic ADR search** Adopt Supabase's built-in pgvector for storing ADR embeddings.	Proposed	AI	PS Priya	2d ago
ADR-004	**Daily AI usage quota at the app layer** Per-user daily AI usage tracked in user_ai_usage.	Accepted	Manual	MC Mira	6d ago
ADR-005	**Prompt injection defence for Search & Ask** Three-layer defence: strict prompt, relevance threshold, moderation.	Accepted	AI · Edited	PS Priya	3d ago
ADR-006	**Monorepo: pnpm workspaces + Turborepo** Three packages: apps/web, apps/api, packages/core.	Accepted	Manual	TR Tomás	Feb 22, 2026
ADR-007	**JSONB storage for Tiptap document trees** ADR content stored as JSONB in adrs.content; rendered from JSON.	Accepted	Manual	JA Jon	Mar 11, 2026
ADR-008	**Toast messaging as default feedback pattern** Non-blocking transient feedback delivered via a global toast system.	Accepted	Manual	MC Mira	Today
ADR-012	**Original HTML-string storage for ADR content** Superseded by ADR-007 in favour of JSONB storage.	Superseded	Manual	TR Tomás	Mar 11, 2026
ADR-015	**OpenAI text-embedding-3-small for ADR embeddings** 1536-dim embeddings generated on ADR save/update.	Draft	AI	PS Priya	Today

ADR-001 ● Accepted Manual

Authentication & authorisation flow

Status ▾ ✏ Edit

✨ AI SUMMARY Regenerate

Session handling and auth mutations (signup, signin, signout, refresh) communicate directly with Supabase Auth from the web client. All other data/business operations route through the Fastify REST API. This keeps auth latency low without compromising the backend-only data rule.

Context

Our initial position required every client call to route through the Fastify REST API. For Supabase Auth specifically, that adds a hop with no security benefit — Supabase Auth is already a trusted identity provider with cookie-based sessions.

Decision

Authentication operations (`signup` , `signin` , `signout` , token refresh) call `@supabase/ssr` directly from the Next.js app. All data and business operations continue to route through Fastify. RLS policies remain the authoritative gate at the database layer.

Consequences

- Lower auth latency; no duplicate session cache between web and API.
- NFR-2.1 clarified in the PRD to scope the “no direct DB” rule to data operations only.
- The Fastify API trusts the Supabase JWT and verifies it on every protected route.

Alternative solutions

> Proxying Supabase Auth through Fastify was considered. It added ~80ms per call and forced us to re-implement cookie refresh semantics. Rejected.

AUTHOR
MC Mira Chen

CREATED
Mar 4, 2026

LAST UPDATED
Apr 2, 2026

CREATION METHOD
Manual

LINKED ADRS Link

DEPENDS ON — Monorepo: pnpm workspaces + Turborepo
ADR-006 · Accepted

INFORMS — Tiptap as the WYSIWYG editor
ADR-002 · Accepted

← All ADRs

✨ AI Brain Dump

Paste rough notes. The assistant drafts a structured ADR you refine before saving.

✓ Brain dump

Unstructured notes — the AI will structure it for you.

we hit p95 of 5s on image thumbs during the campaign last week, origin CPUs pegged. considered varnish but ops pushback. cloudflare seems fine, we already have the account. immutable headers for hashed urls, SWR for mutable. need to document cdn purge as a new ops concern.

312 / 5,000 characters

✨ Generate draft

2 Review & refine

Edit the AI output in the WYSIWYG editor.

✨ AI-generated

Title

Adopt Cloudflare CDN in front of the image service

Status

• Draft • Proposed • Accepted

Content

Heading 2 ▾ | B I | • 1. | </> "

Context

We currently serve image thumbnails from the origin server with no intermediate caching. Peak traffic consistently saturates the image-processing workers, producing 3–5s p95 response times during marketing campaigns.

Decision

Introduce a CDN layer (Cloudflare) in front of the image service. Origin responses will carry `Cache-Control: public, max-age=31536000, immutable` for hashed URLs; mutable URLs will use `stale-while-revalidate`.

Consequences

- Drops origin bandwidth by an estimated 85% based on current hit-ratio projections.
- Introduces CDN purge as a new ops concern; documented in the runbook.

2,384 / 15,000 characters

Daily AI usage 19%
40,310 tokens left · resets 00:00 UTC

MC Mira Chen
mira@acme.dev

Search & Ask

Ask in plain language. Answers are grounded in your ADRs with citations.

Why did we pick pgvector over a dedicated vector database?

Why Tiptap over Lexical? | How do we prevent prompt injection? | Where is ADR content stored?

⌘ Enter to send

ANSWER

Three layers of defence. First, a strict system prompt fixes the assistant's persona and instructs it to refuse any request outside answering from ADR context. Second, before the model is called, the query is vectorised and matched against the corpus — if the best cosine similarity is below 0.78 the system short-circuits with a canned "no matching ADR" response. Third, every user input is routed through the OpenAI Moderation API and rejected if it returns a flag.

CITATIONS

1 Prompt injection defence for Search & Ask
ADR-005 · Accepted

2 pgvector for semantic ADR search
ADR-003 · Proposed